The Live-Graph Software Stack & Agentic Protocol: Architectural Specifications for 2026 City-Level Discovery Infrastructure

The transition from traditional search-centric indexing to agentic discovery layers marks a definitive pivot in the local economic infrastructure of 2026. This architectural shift, exemplified by the LocalPlus framework, replaces the antiquated "Keyword-Link-Click" paradigm with a more sophisticated "Intent-Answer-Action" model designed for the generative era.[1] At the heart of this transformation is the "Live-Graph" software stack, a production-grade ecosystem engineered to mitigate the "Hyper-Local Decay" problem—the tendency of big-tech language models to rely on stale data regarding operating hours, stock availability, and civic hierarchies.[1] To sustain a 99% accuracy rate at the city level, the stack must synthesize real-time data streaming, native graph traversals, and multi-objective reranking into a unified agentic protocol.[1]

The 2026 software ecosystem is defined by its ability to process high-frequency updates while maintaining sub-300ms retrieval latencies, a requirement that necessitates a departure from standard vector-only Retrieval-Augmented Generation (RAG).[1] Instead, the stack utilizes the Model Context Protocol (MCP) as a standardized orchestration bridge, ensuring that the language model (LLM) remains grounded in a persistent, two-way connection with local data sources.[4] This report investigates the technical integration of MCP, the performance benchmarks of Graph-RAG versus multi-modal vector stores, the engineering of symbolic injection layers for commercial disclosure, and the scaling of AI observability metrics like "Answer Share" across millions of user sessions.[1]

Model Context Protocol (MCP) as the Primary Orchestration Bridge

The Model Context Protocol (MCP), standardized by industry leaders such as Anthropic and Google in late 2024 and 2025, serves as the "USB-C for AI," providing a universal standard for connecting AI applications to external data, tools, and resources.[4] In the LocalPlus ecosystem, MCP acts as the primary communication bridge between the Partner App—which ingests real-time merchant data—and the core RAG engine.[4] This protocol addresses the "N x M integration problem," where each AI application previously required custom-coded bridges for every unique data source; with MCP, any server and any client implementing the protocol can interact seamlessly without bespoke integration code.[4]

The MCP architecture consists of three primary components: the MCP Host (the AI application

environment), the MCP Client (the bridge embedded within the host), and the MCP Server (the external service providing context or capabilities).[4] For city-level discovery, this modularity allows the discovery engine to connect with diverse servers—ranging from high-speed SQL databases for merchant inventory to specialized civic servers for local government hierarchies—over a standardized JSON-RPC 2.0 transport.[9]

Table 1: Model Context Protocol (MCP) Primitive Architecture

Primitive	Functionality in LocalPlus	Technical Mechanism
Tools	Executable actions (e.g., book_table, check_parking)	JSON-RPC 2.0 schema defining function names and arguments.[5]
Resources	Contextual data sources (e.g., CSV schemas, API responses)	Identified by unique URIs; enables file and database reads.[9]
Prompts	Reusable templates for grounding specific user intents	Standardized templates to jumpstart multi-turn conversations.[5]
Sampling	Server-initiated requests for LLM completions	Allows servers to "ask the LLM's brain" for complex analysis during tool execution.[5]
Notifications	Real-time state-change alerts	One-way push messages without requiring a request-id or response.[5]

The critical question for the 2026 stack is whether MCP can handle "Streaming Notifications" to update the LLM's context window in real-time without the overhead of a full vector re-index.[1] Technical analysis of the protocol’s lifecycle confirms that MCP is uniquely designed for this requirement through its stateful session management.[5] During the initial handshake, the server and client negotiate capabilities; if the server declares "listChanged": true, it reserves the right to push notifications whenever its internal state changes.[5] For instance, if a local business updates its "star" status or changes its available inventory, the MCP server sends a notifications/tools/list_changed or notifications/resources/list_changed message.[5]

The MCP client reacts to these notifications by performing a targeted refresh of the specific

registry affected, rather than re-scanning the entire ecosystem.[5] In 2026 production environments, this is often implemented over a **Server-Sent Events (SSE)** transport, which maintains a hanging HTTP connection that allows the server to stream context updates directly to the client as they occur.[5] This mechanism ensures that the LLM always operates on the most current metadata, effectively grounding the "Answer Engine" in live data without the latency of traditional background indexing tasks.[5] Recent developments in services like Amazon Bedrock AgentCore Memory have further optimized this by supporting streaming notifications for long-term memory, eliminating the need for polling and ensuring that insights from user interactions are instantly available for context conditioning.[14]

High-Frequency Retrieval: Graph-RAG versus Multi-Modal Vector Stores

The retrieval stage of the 2026 "Answer Engine" must navigate the tension between semantic similarity and complex relational logic. For local discovery, a standard query often involves multiple constraints that are inherently relational: proximity to a secondary location, real-time availability status, and reputation-based eligibility.[1] While vector databases like Pinecone and Milvus excel at finding "near matches" for semantic intent, they struggle with "multi-hop" queries that require joining data across different entities.[16]

Neo4j AuraDB (Graph) provides a native graph engine that treats relationships as first-class citizens, storing them as physical pointers on disk—a technique known as "index-free adjacency".[2] This allows for complex traversals, such as moving from a Business node to a Status node and then to a Geographic_Zone node, without the computational cost of SQL joins or the ambiguity of vector similarity.[2] In February 2026, Neo4j launched its **Aura Agent**, an agent-creation platform that integrates knowledge graphs directly into the agent lifecycle.[17] This platform utilizes a ReAct agent loop to enable chain-of-thought reasoning over graph structures, allowing the discovery engine to resolve queries like "Find a café with 'star' status open now near a 'dagger' cited parking lot" with high precision.[17]

Table 2: 2026 Performance and Capability Matrix: Neo4j vs. Vector Stores

Metric	Neo4j AuraDB (Graph)	Pinecone / Milvus (Vector)
Simple Latency (P95)	50-100ms (Internal execution) [3]	10-50ms (Similarity search) [19]
Multi-Hop Query Logic	Native (Index-free	External (Logic-layer joins

	adjacency) [2]	required) [21]
2026 RAG Integration	Aura Agent (Direct Cypher generation) [17]	Vercel AI SDK / LangChain [22]
Consistency Model	ACID-compliant transactions [16]	Eventually consistent [21]
Multi-Modal Support	Via Aura Agent Similarity Search [17]	Native dense/sparse vector indexing [16]
Throughput (QPS)	Scales with "Secondaries" (Read-only copies) [24]	Up to 10,000 QPS (Serverless) [19]

For LocalPlus, minimizing retrieval latency below the 300ms threshold is a primary production requirement.[1] Technical audits of Neo4j driver performance indicate that while internal execution on the AuraDB server may take only ~82ms, overall response time is often inflated by network round-trips for database discovery and connection establishment.[3] To stay within the 300ms budget, production stacks in 2026 utilize **Secondaries**—static, read-only copies of the primary AuraDB instance that can be scaled up to 15 copies per instance to handle read-heavy workloads.[24] Furthermore, by specifying the database name (always 'neo4j' for Aura) in the driver session, engineers can eliminate one unnecessary round-trip for database discovery, a crucial optimization for low-latency agentic flows.[3]

While Neo4j provides superior relational reasoning, vector stores like Pinecone remain relevant for broad semantic filtering. Pinecone's serverless architecture handles billions of vectors with single-digit millisecond latency and supports "sparse-dense hybrid search," which combines keyword exactness with semantic intent.[16] However, for the specific "city-level" discovery problem, the **ArcadeDB** alternative emerged in 2026 as a strong contender, offering a multi-model engine that supports SQL, Cypher, and Vector search natively in a single engine.[25] ArcadeDB's benchmarks in the LDBC Graphalytics suite showed it outperforming Neo4j 2026 in algorithms like PageRank (0.48s vs 11.15s) and Breadth-First Search (0.13s vs 1.91s), suggesting that multi-model engines may be the most efficient path for production-grade stacks requiring both graph traversals and vector similarity within the same transaction.[25]

The Symbolic Injection Layer: Middleware and Post-Retrieval Mapping

To maintain user trust while monetizing discovery, LocalPlus employs a "Symbolic Disclosure System" that embeds interactive status indicators directly into the AI response.[1] These symbols—Middle Dot (·) for promoted mentions, Dagger (†) for priority citations, and Star (★) for premium partners—must be programmatically injected into the streaming LLM response

based on metadata stored in the Knowledge Graph.[1] This requirement creates a complex software pattern for "Post-Retrieval Symbol Mapping" that operates within the streaming pipeline.[1]

The implementation pattern for 2026 utilizes Next.js middleware and the Vercel AI SDK's streamText function.[22] The process involves several distinct engineering layers:

1. **Metadata Interception:** As the RAG engine retrieves Business IDs from the graph, their sponsorship attributes and status symbols are attached to the metadata payload.[1]
2. **Streaming Middleware:** A custom interceptor monitors the incoming token stream from the LLM. When the model generates the name of a business that exists in the current retrieval context, the middleware cross-references the business ID to find corresponding symbols.[26]
3. **Token Augmentation:** The middleware injects the specific UTF-8 symbol character (e.g., \u2022 for the middle dot) immediately following the business name in the text stream.[26]
4. **Interactive Rendering:** To ensure symbols are interactive (Progressive Disclosure), the frontend must handle partial data parsing.[26]

In Next.js 16+ environments, developers use the **partial-json** library to handle incomplete JSON structures in the stream.[26] This allows the React component to begin rendering the "Disclosure Tooltip" as soon as the symbol appears, even if the model is still generating the rest of the sentence.[26] By wrapping the injected symbols in a custom <DisclosureIcon /> component, the system can trigger a hover-state or click-state that reveals the underlying commercial relationship, fulfilling the requirement for progressive disclosure.[29]

Table 3: Progressive Disclosure Design and Middleware Logic

Interaction Stage	Component Behavior	Technical Implementation
Initial Exposure	Symbol appears inline (·, †, ★)	UTF-8 character injection via server-side stream interceptor.[26]
User Hover	Brief explanation of symbol (e.g., "Paid Citation")	React Popover triggered by symbol metadata.[26]
User Click	Full disclosure modal with transparency data	Next.js Server Action fetching deep metadata from Knowledge Graph.[1]
Transactional Arrow	Direct link to	Dynamic rendering of

	booking/purchase API	<ArrowButton /> using insertAdjacentText().[1]

A critical security and performance challenge in this layer is rendering streamed Markdown or rich text without causing layout shifts or "jank".[33] The 2026 standard for high-performance rendering is the use of a **streaming-markdown parser**, which processes incoming chunks individually and appends them to the DOM via the appendChild() method.[33] This prevents the expensive re-parsing of the entire message with each new token and allows for the seamless integration of the symbolic components within the text flow.[33] Furthermore, a DOM sanitizer is integrated into the renderer to catch potentially insecure output mid-stream, ensuring that the "Language of Visibility" cannot be manipulated into executing malicious scripts.[33]

Mathematical Foundations of Multi-Objective Reranking Models

The LocalPlus ranking engine utilizes a weighted multi-objective scoring model to determine the order of answers: 0.45 Relevance, 0.20 Proximity, 0.15 Reputation, 0.10 Sponsorship, and 0.10 Engagement.[1] To technically implement the **Sponsorship Weight (0.10)** into the reranking math without degrading the core **Relevance (0.45)**, the stack employs latest Cross-Encoder rerankers such as **BGE-Reranker-v2-m3** or **Cohere Rerank 3.5**.[34]

Unlike simple embedding models, Cross-Encoders take a query and document pair as a single input and output a raw relevance logit.[34] For business logic, these raw scores must be mapped to a standardized float value in $$ using a **sigmoid function**.[34] Once the base relevance score (S_{rel}) is normalized, the sponsorship logic is integrated via **Linear Scalarization**.[36]

The mathematical implementation for a candidate item i in production is as follows:

$$Score_{total}(i) = (w_{rel} \cdot S_{rel}) + (w_{prox} \cdot S_{prox}) + (w_{rep} \cdot S_{rep}) + (w_{spon} \cdot B_{paid}$$

To prevent "Sponsorship Decay"—where a paid result significantly degrades the relevance of the answer—the system applies **Augmented Lagrangian-based methods**.[37] This constrained optimization framework allows the model to promote items with higher compensation (sponsorship) while ensuring the final list-level utility stays as close as possible to the original relevance-only ranking.[37] In 2026, this is often handled by a **hypernetwork** that generates parameters for the reranking model according to real-time preference weights, allowing the

discovery engine to adapt to different user groups or holiday promotions without retraining.[36]

Table 4: Reranking Performance and Model Comparison (2026)

Model	Architecture	Context Length	Latency (100 docs)	Best For
Cohere Rerank 3.5	Custom Cross-Encoder	4,096 tokens [35]	~0.5s (512 token docs) [35]	Absolute quality in complex English queries.[35]
BGE-Reranker-v2-m3	Cross-Attention	512-8k tokens [40]	~3.0s (High resource) [40]	Strong multilingual and high-efficiency scenarios.[34]
mxbai-rerank-v2	Qwen-2.5 based	8K - 32K tokens [40]	~0.89s [40]	JSON and structured metadata support.[40]
GTE-ModernBERT	ModernBERT base	149M Params [41]	<0.1s [39]	High-throughput, low-latency deployments.[39]

In production, the **WeightedRanker** strategy (standard in Milvus 2026) is the preferred implementation for these multi-objective formulas.[42] It normalizes scores from different vector search paths (e.g., Image Similarity vs. Text Relevance) using an arctan function and then applies the weights.[42] For LocalPlus, this means that even if a business has a maximum sponsorship weight (0.10), it cannot overcome a significantly lower relevance score (0.45), ensuring that the system remains an "Answer Engine" rather than a purely sponsored billboard.[1]

AI Observability: Tracking "Answer Share" and Citation vs. Mention Metrics

The transition from CPC (Cost Per Click) to "Cost Per Answer Share" requires a specialized observability stack to track visibility across 1,000,000+ distributed user sessions.[1] "Answer Share" is the primary KPI, defined as the percentage of AI-generated responses within a

category/location that include a specific business.[1] To track this at scale, the 2026 stack utilizes observability tools like **Maxim AI, Langfuse, or Arize**, integrated with high-performance telemetry platforms like **Conductor or Rankshift**.[6]

In this environment, visibility is bifurcated into two critical metrics:

- **Mentions (Awareness):** When the AI references a brand in the response text without a link.[44] This measures "conversational share of voice".[44]
- **Citations (Authority):** When the AI explicitly cites a brand as a source, typically including a clickable link.[6] This is the primary measure of authoritative content.[44]

The gap between these two—the "Mention-Citation Gap"—is a key strategic indicator.[44] If a brand is frequently mentioned but never cited, it signals a content gap where the AI model knows the brand but does not trust its content enough to use it as a source.[44] Closing this gap is the foundation of **Answer Engine Optimization (AEO)**.[44]

Table 5: Answer Share and Visibility Metric Framework

Metric	Definition / Formula	Business Implication
Visibility Score	$\left(\frac{\text{Responses mentioning Brand}}{\text{Total Responses}}\right)$ [6]	Benchmarks raw brand recall across major AI engines.[6]
Citation Rate	$\left(\frac{\text{Answers citing Brand URL}}{\text{Total Answers with sources}}\right) \times$ [6]	Signals which domains AI models trust to ground their answers.[6]
Answer Share	Bubble chart of Visibility (X) vs. Citation Rate (Y) [6]	Combined score reflecting overall authority and market presence.[6]
Mention Rate	% of sampled answers referencing specific content [46]	Directional indicator of brand awareness in AI prose.[45]

To aggregate these metrics across millions of sessions, the 2026 stack employs **spatial and temporal aggregation** via Prometheus and Recording Rules.[47] Logging millions of individual "Mention" events per minute would overwhelm standard monitoring instances. Smarter aggregation reduces data volume while preserving the signal [47]:

- **L1 Recording Rules:** Aggregate pod-level logs to the service level (e.g., total mentions of BusinessID_123 across all inference nodes).[47]

- **L2 Recording Rules:** Aggregate service-level data to the environment/city level (e.g., total answer share for the "Cafe" category in Melbourne).[47]
- **Spatial Aggregation:** Reduces cardinality by aggregating across labels like Geo_Zone or Business_Category instead of tracking individual user IDs.[47]

This tiered retention strategy (raw 15s data → 5m averages → 1h summaries) allows for the long-term tracking of Answer Share trends without the massive storage costs associated with raw session logs.[47] For high-stakes B2B sales, this data is visualized in **bubble charts**, where the movement of a bubble to the right indicates improved visibility, and an upward shift indicates an increased citation rate.[6]

Regulatory Compliance and Ethical Symbolic Disclosure

The 2026 software ecosystem is governed by stringent transparency requirements, most notably the **EU AI Act**, which mandates that users be informed when they are interacting with an AI system or when content is AI-generated.[48] Article 50 of the Act requires providers to ensure outputs are marked in a "machine-readable format" and detectable as artificially generated.[32]

For LocalPlus, the symbolic injection system must meet the "Clear and Conspicuous" disclosure standard.[49] In 2026, California's SB 942 and New York's "Synthetic Performer" laws impose similar obligations, requiring that disclosures be visible at the "first interaction or exposure".[48] The stack fulfills these requirements through a multi-layered disclosure approach:

1. **Latent Disclosure:** Digitally signed metadata indicating the AI-generated nature of the response is embedded in the JSON-RPC stream.[48]
2. **Manifest Disclosure:** The symbols (·, †, ★) serve as prominent markings.[48]
3. **Adapted Disclaimers:** For space-constrained interfaces (like mobile chat), the system uses the FEC's "adapted disclaimer" pattern, where a symbol acts as an "indicator" that, upon one action (hover/click), reveals the full "paid for by" disclosure through a "mechanism" like a pop-up or hyperlink.[50]

This regulatory-governance framework ensures that the "Language of Visibility" is not merely a design choice but a robust evidentiary system that complies with the engineering and governance mandates of the 2026 legal landscape.[52]

Conclusions and Production Roadmap for City-Level Answer Engines

The architecture of a 2026 "Answer Engine" is defined by the move from passive retrieval to agentic orchestration. The Model Context Protocol (MCP) provides the necessary bridge to

real-time city data, enabling high-frequency context updates without the latency of re-indexing. By prioritizing Neo4j AuraDB's graph traversals for multi-hop queries and utilizing multi-objective reranking with sigmoid-normalized relevance scores, the stack maintains high grounding accuracy even in commercially influenced environments. The symbolic injection layer, managed through Next.js middleware and partial JSON parsing, offers a path for monetized discovery that remains compliant with the EU AI Act's "clear and conspicuous" mandates. As these systems scale, the use of high-cardinality aggregation via Prometheus will be essential to track "Answer Share" as the definitive metric of the generative era. The future of local discovery depends on this "Live-Graph" stack—a system where data decay is minimized, relationships are first-class citizens, and commercial influence is disclosed through a standardized, interactive language of visibility.

Works cited

1. LocalPlus- Project Grounding Pack (v1.0).rtf
2. Neo4j AuraDB: Fully Managed Graph Database, accessed March 22, 2026, https://neo4j.com/product/auradb/
3. Query execution takes significantly more time while using Neo4j drivers, accessed March 22, 2026, https://support.neo4j.com/s/article/16739301184403-Query-execution-takes-significantly-more-time-while-using-Neo4j-drivers
4. What is Model Context Protocol (MCP)? A guide | Google Cloud, accessed March 22, 2026, https://cloud.google.com/discover/what-is-model-context-protocol
5. Architecture overview - Model Context Protocol, accessed March 22, 2026, https://modelcontextprotocol.io/docs/learn/architecture
6. Key AI Search Metrics: What Businesses Need to Track - Rankshift, accessed March 22, 2026, https://www.rankshift.ai/blog/key-ai-search-metrics/
7. What is the Model Context Protocol (MCP)? - Databricks, accessed March 22, 2026, https://www.databricks.com/blog/what-is-model-context-protocol
8. Developer's Guide to AI Agent Protocols, accessed March 22, 2026, https://developers.googleblog.com/developers-guide-to-ai-agent-protocols/
9. Model Context Protocol (MCP) explained: A practical technical overview for developers and architects - CodiLime, accessed March 22, 2026, https://codilime.com/blog/model-context-protocol-explained/
10. What Is the Model Context Protocol (MCP) and How It Works - Descope, accessed March 22, 2026, https://www.descope.com/learn/post/mcp
11. Understanding Model Context Protocol (MCP): Beyond the Hype | by Raj Kundalia | Medium, accessed March 22, 2026, https://medium.com/@rajkundalia/understanding-model-context-protocol-mcp-beyond-the-hype-582ae84d459d
12. model-context-protocol-resources/guides/mcp-client-development-guide.md at main, accessed March 22, 2026, https://github.com/cyanheads/model-context-protocol-resources/blob/main/guides/mcp-client-development-guide.md
13. WebSocket vs SSE: Which One Should You Use?, accessed March 22, 2026,

https://websocket.org/comparisons/sse/
14. Amazon or AWS, accessed March 22, 2026, https://amazonoraws.com/
15. Terraform Feature Lag Tracker Tool | The "Gap of Grief" Audit - Rack2Cloud, accessed March 22, 2026, https://www.rack2cloud.com/terraform-feature-lag-tracker-tool/
16. Neo4j vs Pinecone | Vector Databases Comparison - Keywords AI, accessed March 22, 2026, https://www.keywordsai.co/market-map/compare/neo4j-vs-pinecone
17. Graph-Driven AI for All: Neo4j Aura Agent Enters General Availability, accessed March 22, 2026, https://neo4j.com/blog/agentic-ai/neo4j-launches-aura-agent/
18. What is the most powerful graph database? - Tencent Cloud, accessed March 22, 2026, https://www.tencentcloud.com/techpedia/132711
19. Best Vector Database 2025: Pinecone vs Weaviate vs Qdrant vs Milvus - TensorBlue, accessed March 22, 2026, https://tensorblue.com/blog/vector-database-comparison-pinecone-weaviate-qdrant-milvus-2025
20. Best Open Source Vector Databases 2026 & Comparison - Redis, accessed March 22, 2026, https://redis.io/blog/top-pinecone-alternatives-for-vector-search/
21. Best Vector Databases in 2026: Complete Comparison Guide - Encore, accessed March 22, 2026, https://encore.dev/articles/best-vector-databases
22. Real-time AI in Next.js: How to stream responses with the Vercel AI ..., accessed March 22, 2026, https://blog.logrocket.com/nextjs-vercel-ai-sdk-streaming/
23. Wrote a comparison of open-source Neo4j alternatives in 2026 - the licensing landscape has changed significantly : r/Database - Reddit, accessed March 22, 2026, https://www.reddit.com/r/Database/comments/1rw7qu2/wrote_a_comparison_of_opensource_neo4j/
24. Boosting AuraDB Read Performance With Secondaries - Neo4j, accessed March 22, 2026, https://neo4j.com/blog/auradb/boosting-read-performance-secondaries/
25. Neo4j Alternatives in 2026: A Fair Look at the Open-Source Options - ArcadeDB, accessed March 22, 2026, https://arcadedb.com/blog/neo4j-alternatives-in-2026-a-fair-look-at-the-open-source-options/
26. Next.js + LLMs: The Guide to Streaming Structured Responses - Calfus, accessed March 22, 2026, https://www.calfus.com/post/next-js-llms-the-guide-to-streaming-structured-responses
27. MCP vs API: The Future of Software Integration in AI-Driven Systems - Koombea AI, accessed March 22, 2026, https://ai.koombea.com/blog/mcp-vs-api
28. Next.js Middleware Explained: How It Works, Why It Matters, and Real Production Use Cases | by Rohit Arya, accessed March 22, 2026, https://rohitarya18.medium.com/next-js-middleware-explained-how-it-works-why-it-matters-and-real-production-use-cases-e6e544910868
29. MCP Apps - Model Context Protocol, accessed March 22, 2026, https://modelcontextprotocol.io/extensions/apps/overview

30. Tesslate/Next.js-Dataset - Hugging Face, accessed March 22, 2026, https://huggingface.co/datasets/Tesslate/Next.js-Dataset/viewer/default/train
31. Improve subagent infrastructure and documentation strategy - Amp Code, accessed March 22, 2026, https://ampcode.com/threads/T-391b475d-70ef-43b7-8e14-069d262b3a07?q=librarian
32. Illuminating AI: The EU's First Draft Code of Practice on Transparency for AI-Generated Content | Publications | Kirkland & Ellis LLP, accessed March 22, 2026, https://www.kirkland.com/publications/kirkland-alert/2026/02/illuminating-ai-the-eus-first-draft-code-of-practice-on-transparency-for-ai
33. Best practices to render streamed LLM responses | AI on Chrome, accessed March 22, 2026, https://developer.chrome.com/docs/ai/render-llm-responses
34. BAAI/bge-reranker-v2-m3 · Hugging Face, accessed March 22, 2026, https://huggingface.co/BAAI/bge-reranker-v2-m3
35. Cohere Rerank 3.5 - Oracle Help Center, accessed March 22, 2026, https://docs.oracle.com/en-us/iaas/Content/generative-ai/benchmark-cohere-rerank-3-5.htm
36. Controllable Multi-Objective Re-ranking with Policy Hypernetworks, accessed March 22, 2026, http://ai.ruc.edu.cn/uploads/20230809/1a81c16f7cf590ba0d7a03cb3945abf3.pdf
37. Multi-objective Relevance Ranking - SIGIR eCom, accessed March 22, 2026, https://sigir-ecom.github.io/ecom2019/ecom19Papers/paper30.pdf
38. A Multi-Objective Learning to re-Rank Approach to Optimize Online Marketplaces for Multiple Stakeholders - ResearchGate, accessed March 22, 2026, https://www.researchgate.net/publication/318868520_A_Multi-Objective_Learning_to_re-Rank_Approach_to_Optimize_Online_Marketplaces_for_Multiple_Stakeholders
39. Top 8 Rerankers: Quality vs Cost. A practical benchmark of modern… | by Bhagya Rana | Medium, accessed March 22, 2026, https://medium.com/@bhagyarana80/top-8-rerankers-quality-vs-cost-4e9e63b73de8
40. Baked-in Brilliance: Reranking Meets RL with mxbai-rerank-v2 - Mixedbread, accessed March 22, 2026, https://www.mixedbread.com/blog/mxbai-rerank-v2
41. Reranker Benchmark: Top 8 Models Compared - AIMultiple, accessed March 22, 2026, https://aimultiple.com/rerankers
42. Reranking | Milvus Documentation, accessed March 22, 2026, https://milvus.io/docs/reranking.md
43. RRF Ranker | Cloud | Zilliz Cloud Developer Hub, accessed March 22, 2026, https://docs.zilliz.com/docs/reranking-rrf
44. Mention & Citation Tracking for AI Visibility | Conductor Features, accessed March 22, 2026, https://www.conductor.com/platform/features/ai-search-performance/ai-mention-citation-tracking/
45. How to measure your brand's AI search performance: A complete guide to AI

metrics & KPIs, accessed March 22, 2026, https://envisionitagency.com/blog/ai-performance-metrics-guide/
46. How to Track Your AI Search Visibility: Tools, Metrics, and Provenance - Digital Marketing, accessed March 22, 2026, https://www.fastfrigate.com/resources/ai-search-visibility-tracking-optimization/
47. How to Build Metric Aggregation Strategies - OneUptime, accessed March 22, 2026, https://oneuptime.com/blog/post/2026-01-30-metric-aggregation-strategies/view
48. Article 50: Transparency Obligations for Providers and Deployers of ..., accessed March 22, 2026, https://artificialintelligenceact.eu/article/50/
49. Advertising Law Compliance in 2026: Five Developments Every Advertiser Should Know, accessed March 22, 2026, https://www.afslaw.com/perspectives/the-fine-print/advertising-law-compliance-2026-five-developments-every-advertiser
50. What You Need to Know about the FEC's New Internet Communications Disclaimer Rules, accessed March 22, 2026, https://www.cov.com/en/news-and-insights/insights/2023/02/what-you-need-to-know-about-the-fecs-new-internet-communications-disclaimer-rules
51. AI Trends for 2026 – Return of the Brussels Effect: AI Transparency Requirements Come to California | MoFo Tech, accessed March 22, 2026, https://mofotech.mofo.com/topics/ai-trends-for-2026-return-of-the-brussels-effect-ai-transparency-requirements-come-to-california
52. Transparency obligations for AI-generated content under the EU AI Act: From principle to practice | Herbert Smith Freehills Kramer | Global law firm, accessed March 22, 2026, https://www.hsfkramer.com/notes/ip/2026-03/transparency-obligations-for-ai-generated-content-under-the-eu-ai-act-from-principle-to-practice